



# Machine Learning Project

Mac Lambert  
Self-Paced Data Science  
Flatiron School


# Predicting Customer Churn for SyriaTel

Question: How do we predict whether a customer will churn (or stop doing business) with us?

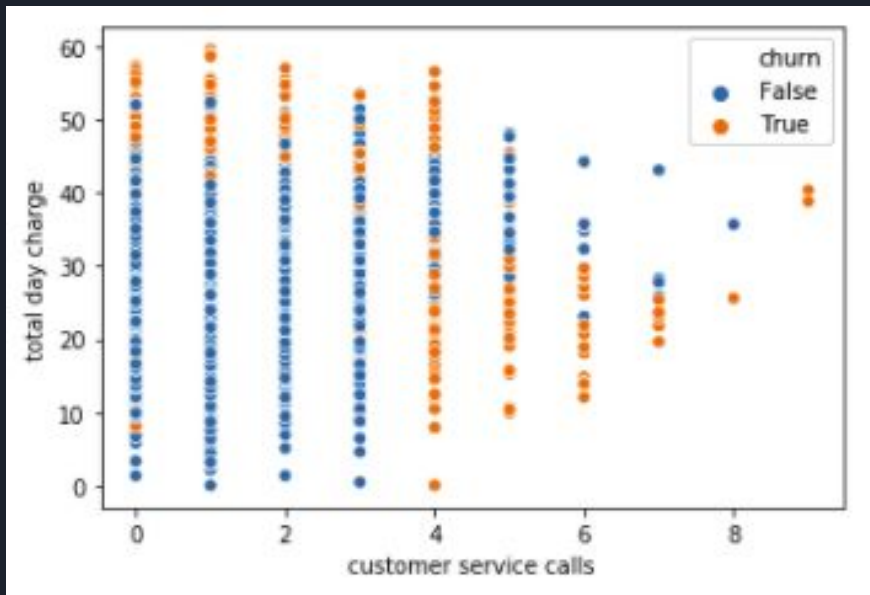
Roughly 15% of customers churn in any given month.

At this rate, over 80% of current customers will be gone within a year!

Customers vs. Month



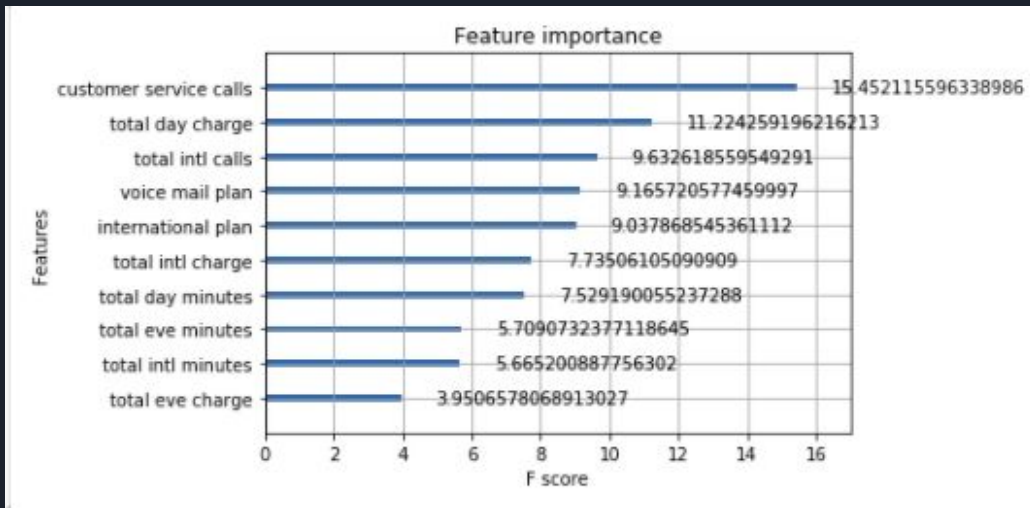
# What do we do about this?



## Methodology:

- Use existing customer info to determine what makes a customer more likely to leave
- Our machine learning model built with customer info predicts churn to 97% accuracy

# What should we focus on?



- The highest ranked features are shown on the left.
- Recommendation: focus on customer service calls, total day charge.



## Further Work

- The next step in this analysis would be to examine charges for customers likely to churn, to determine cost savings impact. This is outside the scope of the predictive model discussed.
- Similarly, optimization was only run on the best performing model. Although 97% accuracy is fantastic, we could have spent the time to optimize all models.



# Thank you & Appendix

- Helpful article on interpreting f score of XGBoost:  
<https://towardsdatascience.com/be-careful-when-interpreting-your-features-importance-in-xgboost-6e16132588e7>
- Link to Kaggle dataset: <https://www.kaggle.com/becksddef/churn-in-telecoms-dataset>
- Steps for one hot encoding state information:  
<https://towardsdatascience.com/choosing-the-right-encoding-method-label-vs-onehot-encoder-a4434493149b>